Reinforcement Learning Approach on the Order Scheduling for the RoboCup Logistic league

by Khodachenko Ian

1 Robo Cup Logistic League Motivation

The Robo Cup Logistics League (RCLL) is designed to represent a common scenario in Industry 4.0. Here one needs to assemble products based on dynamic custom orders in a real-time environment. This means that the environment is driven by a consistent stream of orders and individual delivery windows, which dictate the steps robots in the environment need to take. The challenge here consists in scheduling the assembly process of orders in a feasible manner, to abide the delivery window of as many orders as possible, to figure out priorities or complexities and to understand which orders can't be completed in time.

2 Reinforcement Learning approach

We present a solution based on Reinforcement Learning (RL) to tackle the problem of deciding which assembly step the robots need to process next. Due to the complexity of the problem a Deep Q-Learning (DQL) approach is initially applied, while keeping open options and alternatives described in section 2.5.

The goal of Q-learning is for an agent to select actions, which maximize the total reward earned over an episode (=game), by observing the state and reward from the environment (see figure 1). So from a state s and action a we want to get the maximum future reward r, which we denote as the Q-value $Q(s,a) = r(s,a) + \gamma max_a Q(s',a)$. With hyper parameter γ we can scale the future rewards.

To shift from Q-learning to DQL we apply a Deep Neural Network (DNN or DQN) to learn the Q-values for each individual action, given a state (see figure 2). We can update the DNN via stochastic gradient descent and backpropagation from a squared error loss on the target $Q_\theta(s,a)$. The θ is the policy we want to optimize. Since in contrast to a normal DNN scenario, here we have to update the target we are learning (like cutting the branch one sits on), we need to set-up two parallel networks like in figure 3. These networks are of identical architecture, where we want to fixate the target network weights, while training the prediction network. The weights are copied to the target network every C iterations and enable us to form a more stable policy θ. As an addition to this, we apply experience replay, by storing a number of $(s_t, a_t, r(s_t, a_t), s_{t+1})$ tuples, which we can reuse reducing total data needed.

Putting all together we arrive at figure 4, which roughly translates to

1. Giving the current state to the DQN, which returns the Q-values for all possible actions.
2. Selection a possible action with an ϵ-greedy policy, where we select a random action with probability $1 - \epsilon$ or the action with the maximum Q-value $max_a Q(s,a)$.
3. Proceed to next state taking the selected action and receiving respective reward. Also storing the transition in the replay memory.
4. Sample random batches of transitions from the replay buffer to compute the loss $L = (r + \gamma max_{a'} Q_{\theta'}(s',a') - Q_\theta(s,a))^2$.
5. Perform gradient descent minimizing the loss for the prediction network.
6. Every C iteration copy the weights of the prediction network to the target network.
7. Repeat all of the above for a number of episodes.

2.1 Environment

In the first steps we apply a more simplified environment, which only provides an order schedule with delivery windows. The initial goal is thus for the system to learn which orders are most feasible to complete or award most rewards.

We build the environment so that we can have a precision of seconds, while not specifically modelling each
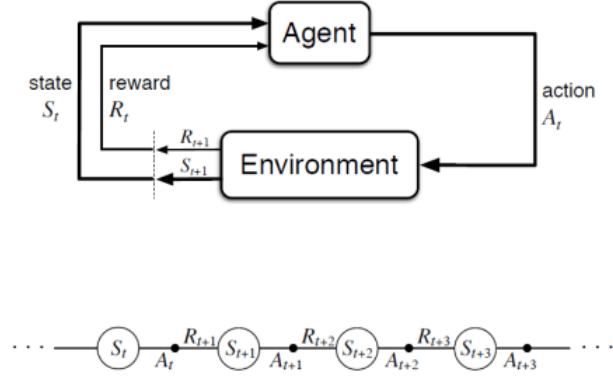
Fig. 1. Basic idea of Reinforcement Learning. Image from [Cho18].
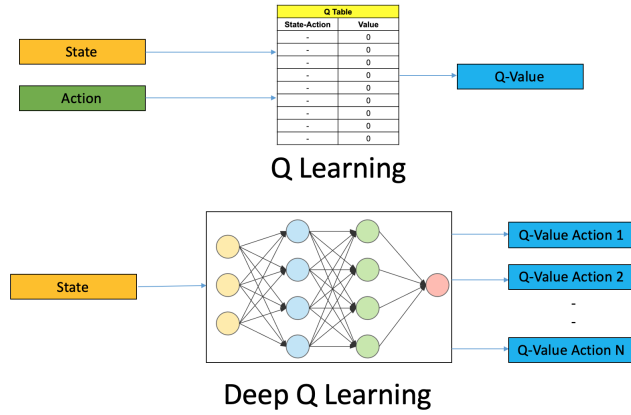
Fig. 2. Comparison of using a Q-table and a neural network. Image from [Cho19].

second as a discrete time step. The idea is to merge the discrete time into the actual state as a numerical feature, which is incremented in the expected passed time for a processing step. Initially the elapsed time for the intermediate steps is drawn from Gaussian distributions, which can be later extended with application of the distance matrix and real-world data.

Each task t_k will represent the duration for a sub-step needed to complete order $O = \sum_i^N t_i$ and will be modelled as $t_k = \mathcal{N}(\mu_k, \sigma_k^2) + \mathcal{U}(a_M, b_M)$. Here the normal distribution $\mathcal{N}$ is given by $\mu_k = a_{ij}$ drawn from a distance matrix representing the distance between machines M_i and M_j. The σ_k can either be fine-tuned by hand or matched from real recorded robot data. In other words $\mathcal{N}$ estimates the time a robot will take to move. For the uniform distribution $\mathcal{U}$ we adapt the values given from the RCLL rulebook, which represent random uniform processing time from a_M to b_M seconds at the specific machine M.

The environment will handle the internal states, which include the playing field distances, product schedule, and robot locations. It will support OpenAI-Gym-like ([BCP$^+$16]) functions and more optimizations. The environment can handle communication with the Referee Box (RefBox) for the RCLL, to pull out the game field, orders and intermediate points as rewards. The reason for this is to learn the specific official implementation and maximize the rewards based on the official tool, as there may be inconsistencies to the rulebook and the environment will need no updates on rule changes (as those are implemented in the RefBox by the official parties).
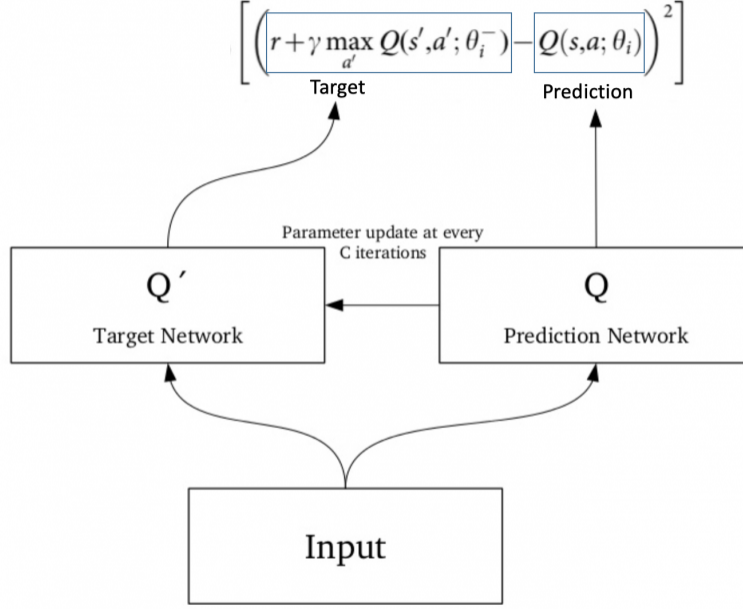
Fig. 3. Structure for updating the two internal networks in the RL set-up. Image from [Cho19].

2.2 States

The state space consists of up to 10 orders (incl. over time), up to 3 products in the pipeline (which corresponds to the physical object one of the 3 robots can transport), distance matrix of the machines, availability of machines and current time. While the distance matrix is constant throughout an episode, the other parameters are simulating a dynamic real-time environment.

Specifically we can represent an order as a vector of $O = \{\mathcal{B}, \mathcal{R}_1, \mathcal{R}_2, \mathcal{R}_3, \mathcal{C}, n, c, d_{start}, d_{end}\}$, where

- $\mathcal{B} = \{red, black, silver\}$ represents a set of base types
- $\mathcal{R} = \{blue, green, orange, yellow, F\}$ is one of up to 3 ring types
- $\mathcal{C} = \{black, grey\}$ is the cap type
- $n = \{1, 2\}$ is the requested amount or products requested
- $c = \{T, F\}$ indicates whether the order is competitive
- $d_{start} = \{d_{start} \in \mathbb{N} | \ 0 \le d_{start} \le 1020\}$ is the start of the delivery window
- $d_{end} = \{d_{end} \in \mathbb{N} | \ 0 \le d_{start} < d_{end} \le 1020\}$ is the end of the delivery window

An intermediate product in the pipeline will look like $P = \{\mathcal{B}, \mathcal{R}_1, \mathcal{R}_2, \mathcal{R}_3\}$. We can also omit the cap as we will deliver straight after. The symmetric distance matrix ($D \in \mathbb{R}^6$) is computed at the start of the episode. The machine availability ($A \in \mathbb{R}^6$) exists mainly for handling real-time events. Finally we also add the running time in seconds as $t = \{t \in \mathbb{N} | \ 0 \le t \le 1020\}$.

While there is still room for changes the current dimensions are

$$10 \cdot O \times 3 \cdot P \times D \times A \times t'$$
$$\approx 2 \cdot 10^4 \cdot (\mathcal{B} \times \mathcal{R} \times \mathcal{C} \times n \times c \times d'_{start} \times d'_{end} \times \mathcal{B} \times \mathcal{R} \times t')$$
$$\approx 35 \cdot 10^6 \cdot (d_{start} \times d_{end} \times t')$$

Noteworthy here is that we can bind $(d_{start} \times d_{end} \times t') \lesssim 22$ because the fact that we will only have increases of about 45 seconds and will only ever see that many combinations.

2.3 Actions

The actions consist of processing (or "getting next") one of the following $A = \{\mathcal{B}_P, \mathcal{R}_P, \mathcal{C}_P, d_P\}$, which is a base, ring, cap or discarding the product in pipeline P. By discarding we offer the option to discontinue the

Fig. 4. Structure for the full DQL cycle. Image from [Cho19].

current order and use the pipeline slot for another order . In total we have an action space of $A \cdot P = 30$, if we let the agent choose from all possible actions. But here we don't want it to guess on possible future orders (it is also not in the RCLL rules), so we can strictly only let it choose actions which would advance one of the current objects in the pipeline. This reduces the maximum number of valid actions we can have at one state to 13 as we can have up to 10 orders and can choose to discard any of the 3 pipelined products.

2.4 Rewards

The initial idea for the rewards is that they should strongly mirror the current rewards assigned by the RefBox, to optimize the points gained in an actual game using the RefBox's rulebook implementation. There are just minor additions to the rewards, like e.g. adding some for acquiring bases to stimulate the learning of the network. This point may see more changes later in the development process, as it is one of the main ways to regulate the learned policy.

2.5 Backup solutions and thoughs

For DQL to work we make sure the Markov property holds, i.e., that each state only depends on the previous state and the transition from that state to the current one. For this reason we need to make sure that the states provide all the needed information, so the above design is not completely fixed and may be updated later.
As such here one will find possible fallback strategies and properties in no particular order:

- We avoid using a multi-agent model by focusing on the order schedule and deducting a task sequence based on it. This task sequence can be forwarded from the GRIPS Teamserver to the individual robots and the time they take will be modelled as described in section ??.
- If we need to reduce complexity or in contrast intertwine the actual robot movement deeper into the model, we can apply similar multi-skill and multi-machine approaches as described in [QWGL16].
- In the case where the estimation of time consumption for the specific tasks do not represent the reality well enough, we can train an additional neural network in a supervised manner. Here the inputs would be from the measured data for a given machine distance matrix - to give us more accurate estimations for the expected time consumption at a specific task.
- When the communication with the RCLL RefereeBox, becomes problematic the rules can be manually translated into the environment.
- Initially we will allow the model to only take actions for known orders (similar to Google's AlphaGo) and thus preventing guessing the future, but this can also be changed by assigning negative rewards. Still the rulebook seem to not support (intermediate) rewards for yet not issued orders.
- There is a possibility that the model will possibly disregard the delivery windows as (per official reward structure) the assembly of more difficult products yields much more then the timely delivery.

– The topic is generally covered with use of RL in a number of papers and to not go into the details of further alternatives: Sadly many use more simple and abstract environments, but the current design takes a couple ideas from them and may adjust more based on [GTGS12], [SLS18], [WRB+18], [WXZL11], [SAM17], [SJ15], [LC15] and some more. Note that the last few are not about RL any more but have good ideas on designing the problem.

3   Current development schedule

First we try to build a network which can assemble a single order by selecting from the actions. Next we want it to recognize and understand the time it will take to process a single order, with a constant distance between machines. To do this we want it to be able to select from a number of orders and try to understand the delivery window.

The environment will be extended as needed and the reward and action space will be adjusted as well. This should enable scaling up the difficulty and observe the feasibility of the DQN approach.

References

BCP+16. Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. Openai gym. https://github.com/openai/gym, 2016.

Cho18.   Ankit Choudhary. Nuts & bolts of reinforcement learning: Model based planning using dynamic programming, 2018.

Cho19.   Ankit Choudhary. A hands-on introduction to deep q-learning using openai gym in python, 2019.

GTGS12. Robert Glaubius, Terry Tidwell, Christopher Gill, and William D Smart. Real-time scheduling via reinforcement learning. *arXiv preprint arXiv:1203.3481*, 2012.

LC15.   James T Lin and Chien-Ming Chen. Simulation optimization approach for hybrid flow shop scheduling problem in semiconductor back-end manufacturing. *Simulation Modelling Practice and Theory*, 51:100–114, 2015.

QWGL16. Shuhui Qu, Jie Wang, Shivani Govil, and James O Leckie. Optimized adaptive scheduling of a manufacturing process system with multi-skill workforce and multiple machine types: An ontology-based, multi-agent reinforcement learning approach. *Procedia CIRP*, 57:55–60, 2016.

SAM17.   Jamal Shahrabi, Mohammad Amin Adibi, and Masoud Mahootchi. A reinforcement learning approach to parameter estimation in dynamic job shop scheduling. *Computers & Industrial Engineering*, 110:75–82, 2017.

SJ15.   Pankaj Sharma and Ajai Jain. Performance analysis of dispatching rules in a stochastic dynamic job shop manufacturing system with sequence-dependent setup times: Simulation approach. *CIRP Journal of Manufacturing Science and Technology*, 10:110–119, 2015.

SLS18.   Yeou-Ren Shiue, Ken-Chuan Lee, and Chao-Ton Su. Real-time scheduling for a smart factory using a reinforcement learning approach. *Computers & Industrial Engineering*, 125:604–614, 2018.

WRB+18. Bernd Waschneck, André Reichstaller, Lenz Belzner, Thomas Altenmüller, Thomas Bauernhansl, Alexander Knapp, and Andreas Kyek. Optimization of global production scheduling with deep reinforcement learning. *Procedia CIRP*, 72(1):1264–1269, 2018.

WXZL11. Jun Wu, Xin Xu, Pengcheng Zhang, and Chunming Liu. A novel multi-agent reinforcement learning approach for job scheduling in grid computing. *Future Generation Computer Systems*, 27(5):430–439, 2011.